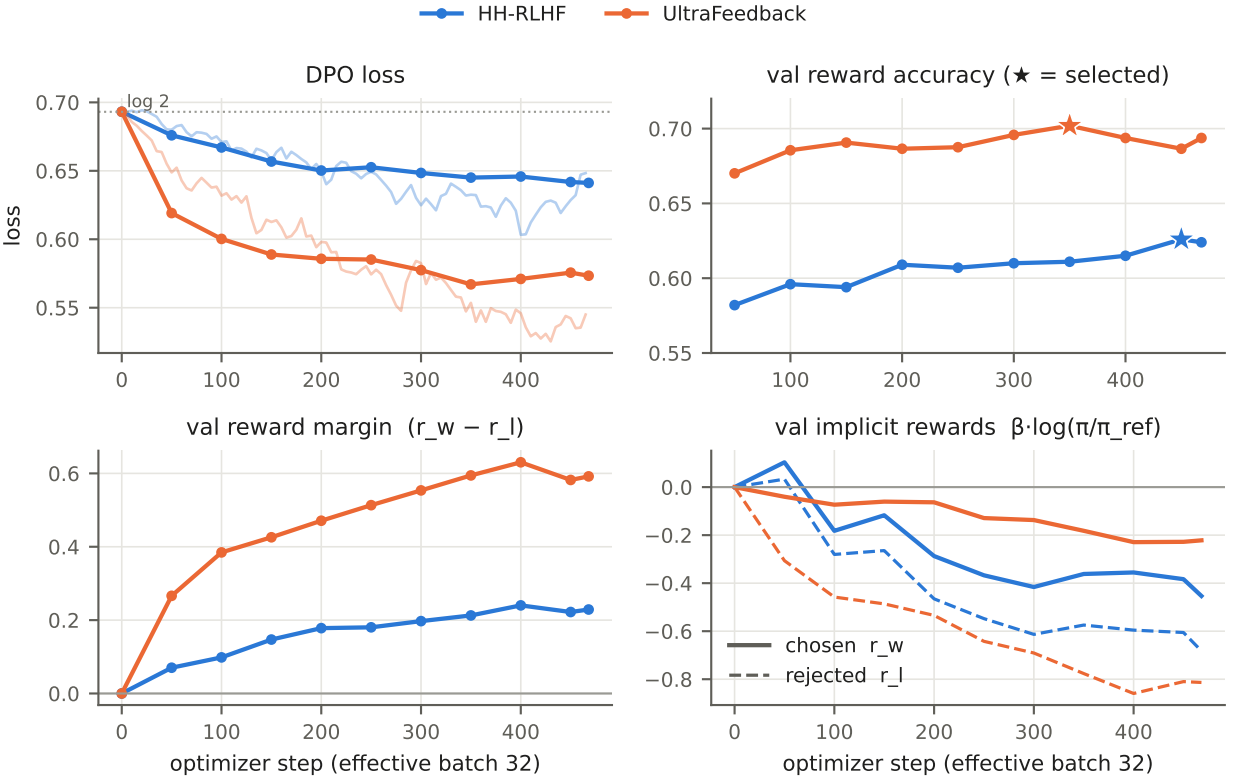
HH-RLHF
UltraFeedback
DPO loss
log 2
loss
0.70
0.65
0.60
0.55
0
100
200
300
400
val reward accuracy (★ = selected)
0.70
0.65
0.60
0.55
100
200
300
400
val reward margin (r_w − r_l)
0.6
0.4
0.2
0.0
0
100
200
300
400
optimizer step (effective batch 32)
val implicit rewards β·log(π/π_ref)
0.0
−0.2
−0.4
−0.6
−0.8
chosen r_w
rejected r_l
0
100
200
300
400
optimizer step (effective batch 32)