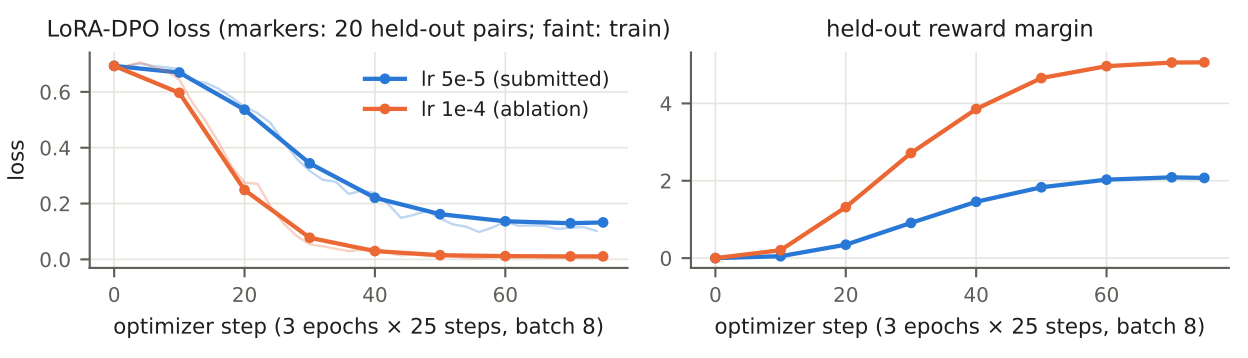
LoRA-DPO loss (markers: 20 held-out pairs; faint: train)
loss
0.0
0.2
0.4
0.6
0
20
40
60
optimizer step (3 epochs × 25 steps, batch 8)
lr 5e-5 (submitted)
lr 1e-4 (ablation)
held-out reward margin
0
2
4
0
20
40
60
optimizer step (3 epochs × 25 steps, batch 8)