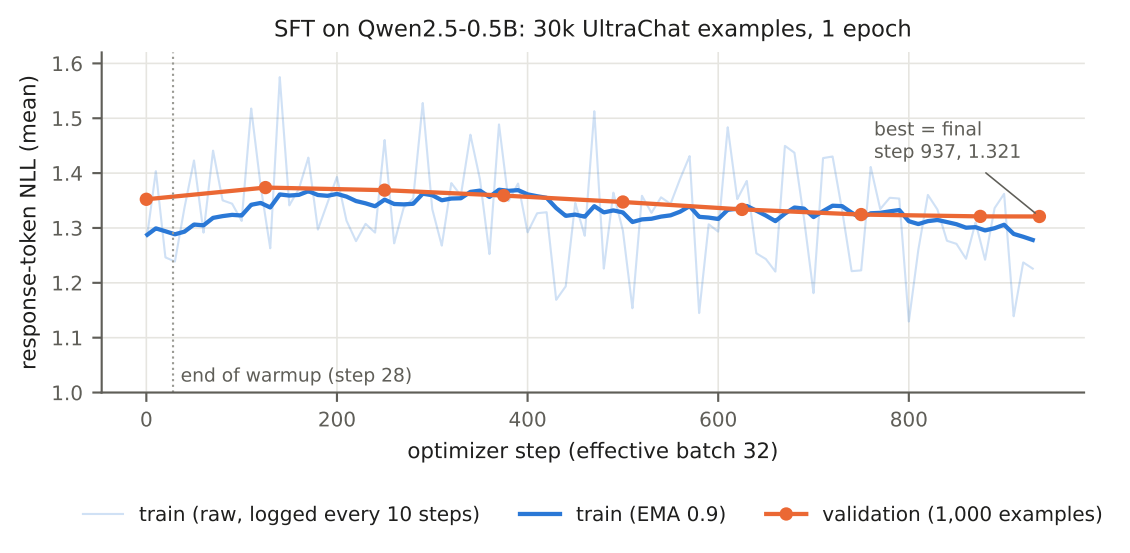
SFT on Qwen2.5-0.5B: 30k UltraChat examples, 1 epoch
response-token NLL (mean)
1.6
1.5
1.4
1.3
1.2
1.1
1.0
best = final
step 937, 1.321
end of warmup (step 28)
0
200
400
600
800
optimizer step (effective batch 32)
train (raw, logged every 10 steps)
train (EMA 0.9)
validation (1,000 examples)